

Minimum Spanning Trees

Spanning subgraph

- Subgraph of a graph G containing all the vertices of G

Spanning tree

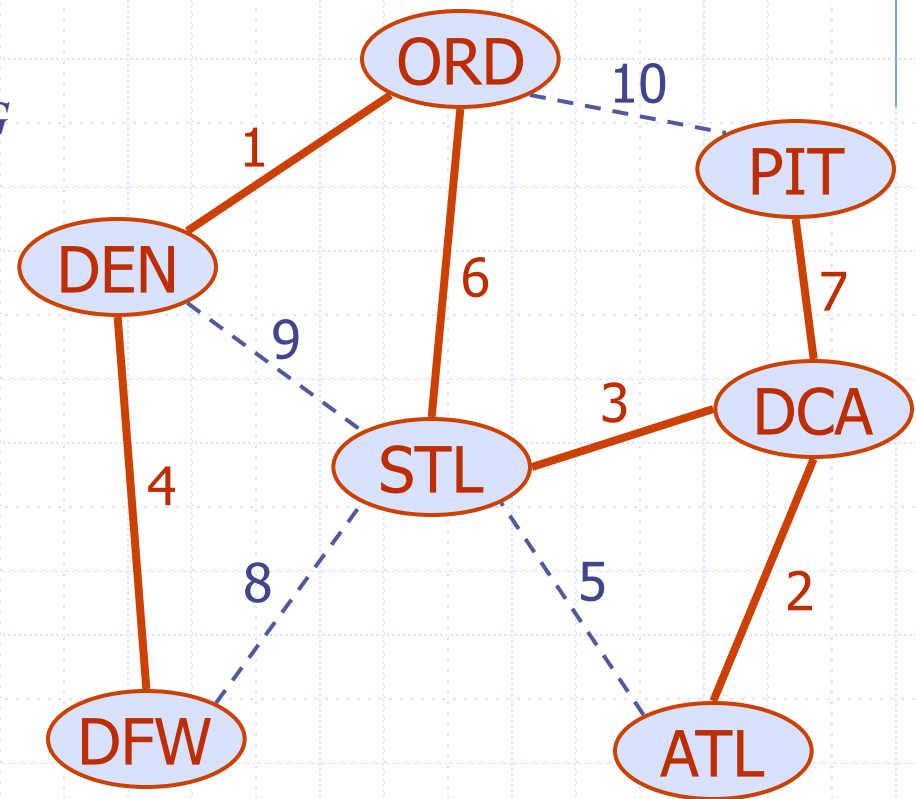
- Spanning subgraph that is itself a (free) tree

Minimum spanning tree (MST)

- Spanning tree of a weighted graph with minimum total edge weight

Applications

- Communications networks
- Transportation networks



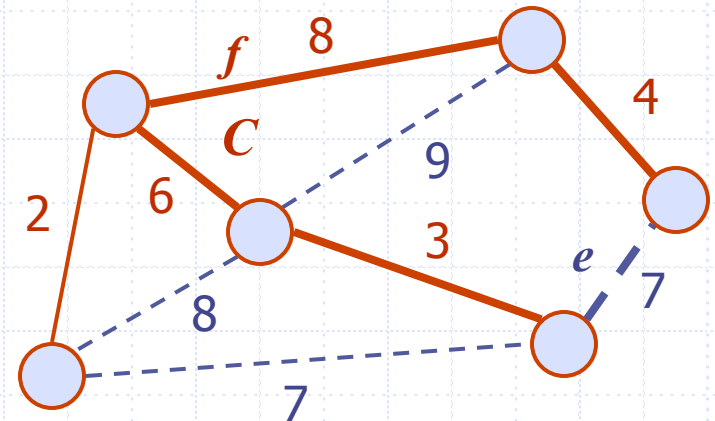
Cycle Property

Cycle Property:

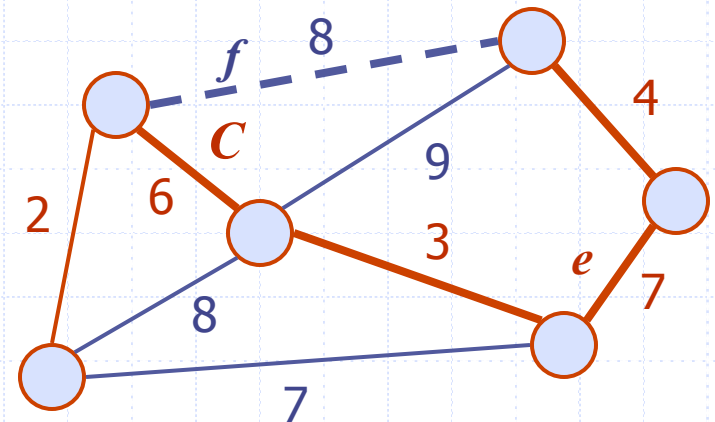
- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $weight(f) \leq weight(e)$

Proof:

- By contradiction
- If $weight(f) > weight(e)$ we can get a spanning tree of smaller weight by replacing e with f



Replacing f with e yields a better spanning tree



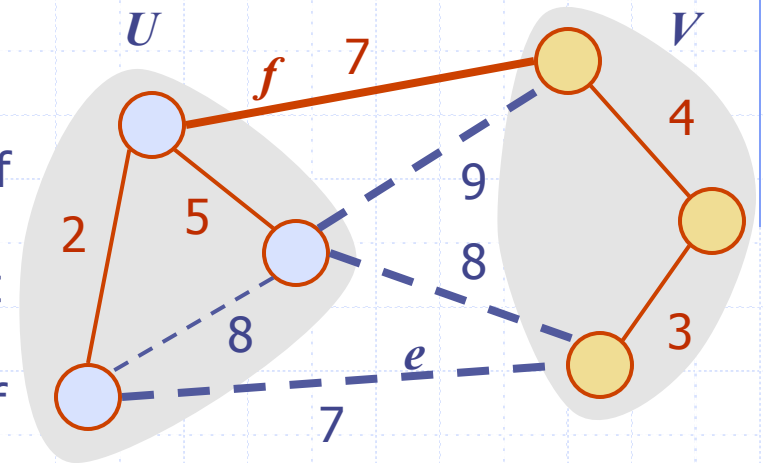
Partition Property

Partition Property:

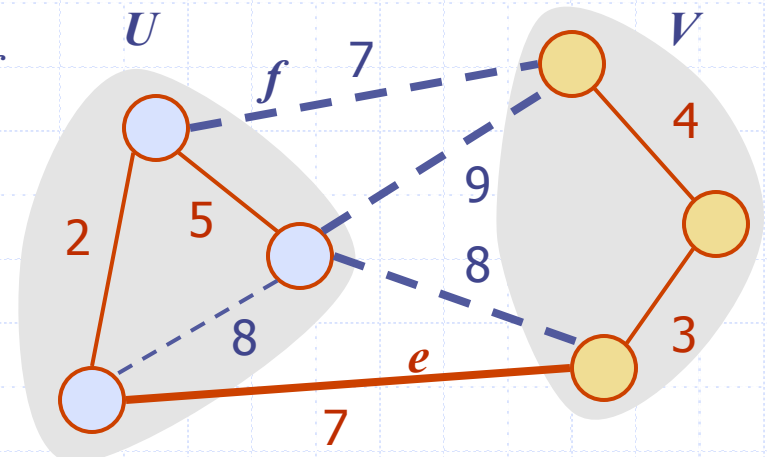
- Consider a partition of the vertices of G into subsets U and V
- Let e be an edge of minimum weight across the partition
- There is a minimum spanning tree of G containing edge e

Proof:

- Let T be an MST of G
- If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition
- By the cycle property,
$$\text{weight}(f) \leq \text{weight}(e)$$
- Thus, $\text{weight}(f) = \text{weight}(e)$
- We obtain another MST by replacing f with e



Replacing f with e yields another MST

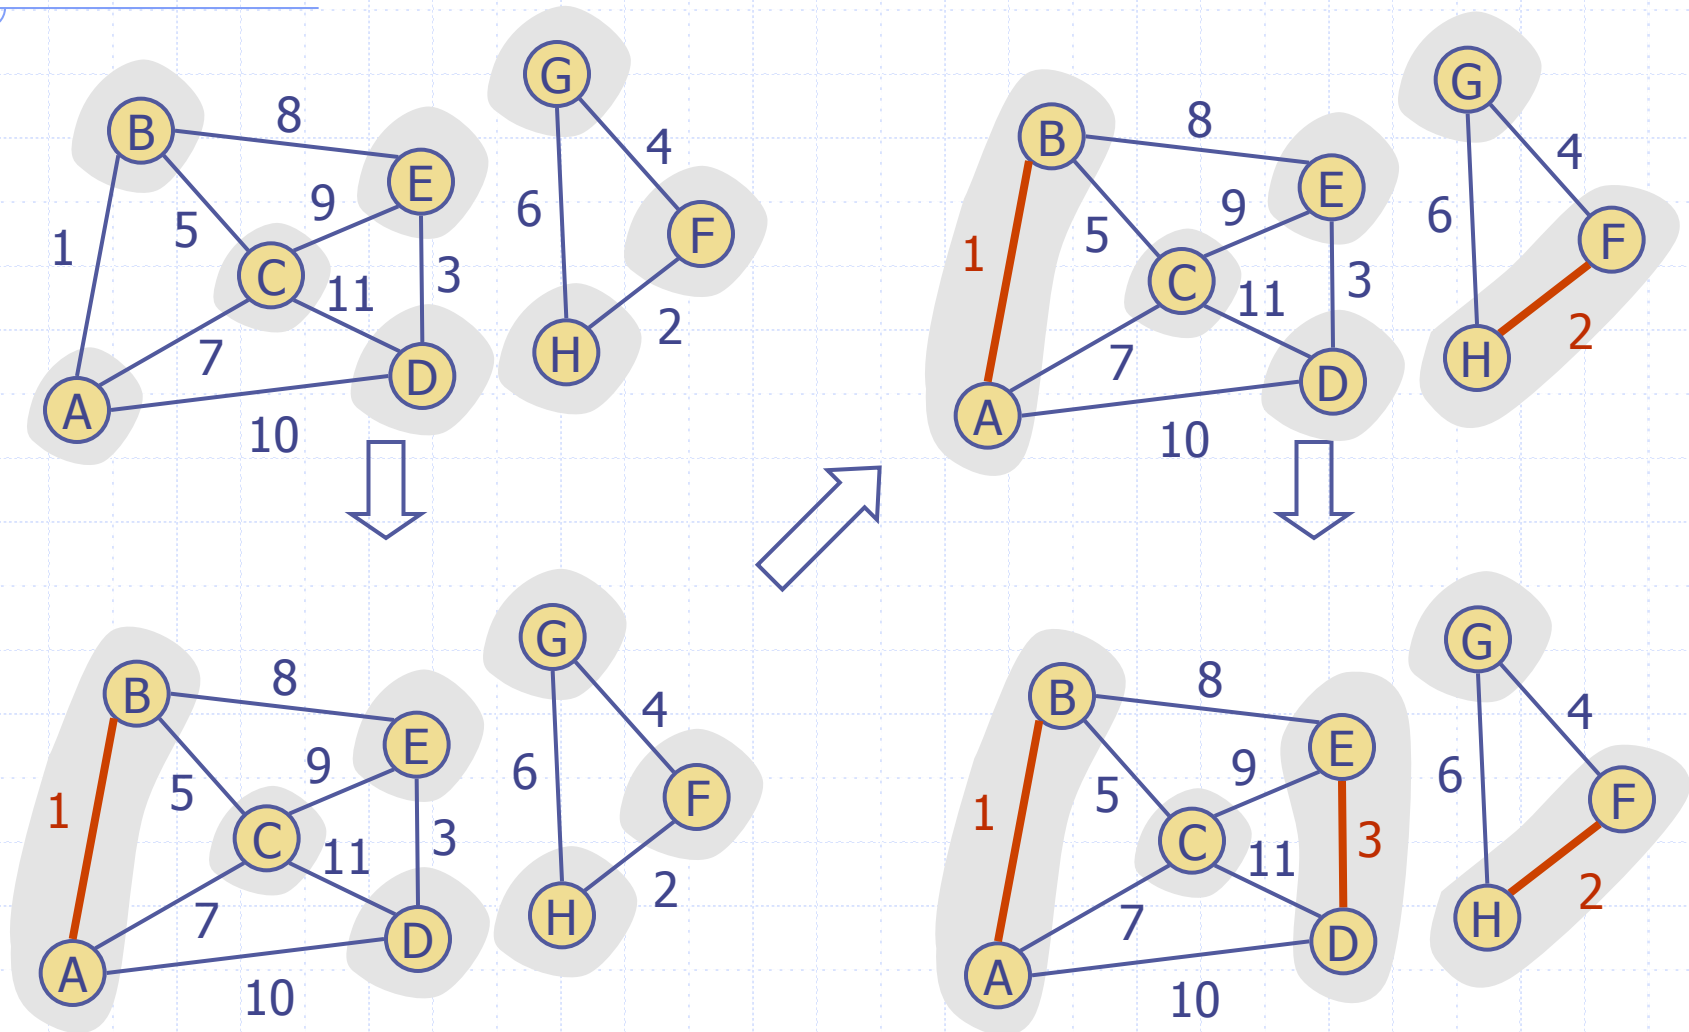


Kruskal's Algorithm

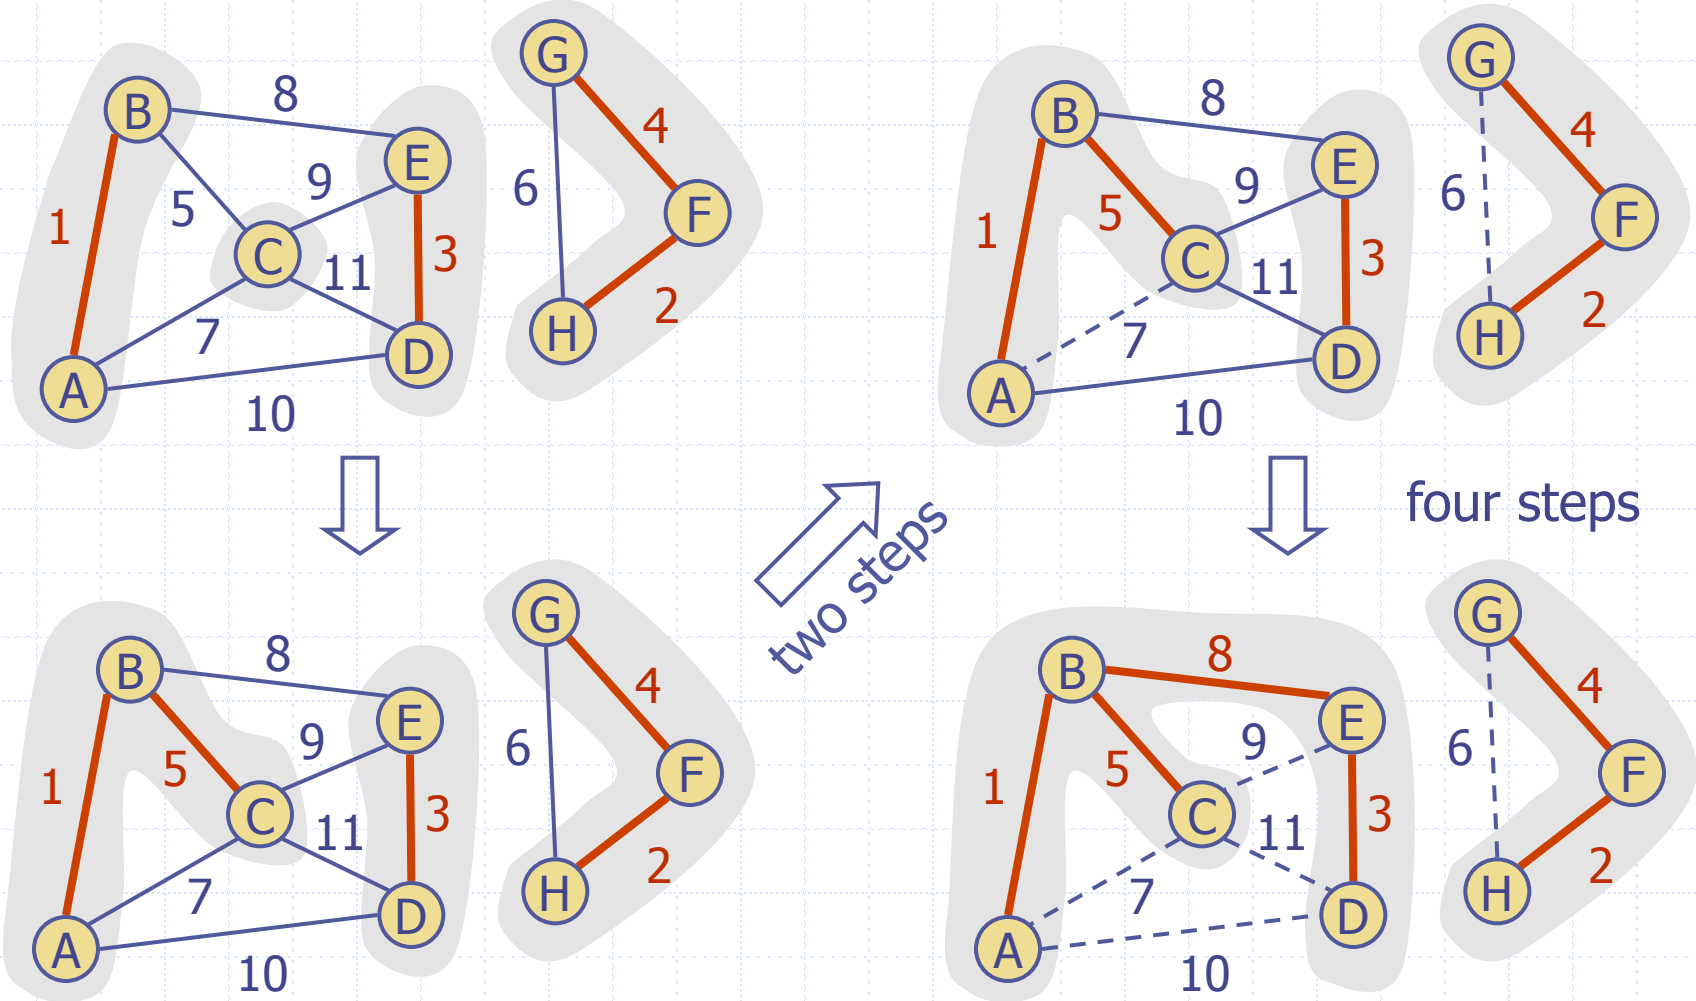
- Maintain a partition of the vertices into clusters
 - Initially, single-vertex clusters
 - Keep an MST for each cluster
 - Merge “closest” clusters and their MSTs
- A priority queue stores the edges outside clusters
 - Key: weight
 - Element: edge
- At the end of the algorithm
 - One cluster and one MST

```
Algorithm KruskalMST( $G$ )  
  for each vertex  $v$  in  $G$  do  
    Create a cluster consisting of  $v$   
  let  $Q$  be a priority queue.  
  Insert all edges into  $Q$   
   $T \leftarrow \emptyset$   
  { $T$  is the union of the MSTs of the clusters}  
  while  $T$  has fewer than  $n - 1$  edges do  
     $e \leftarrow Q.removeMin().getValue()$   
     $[u, v] \leftarrow G.endVertices(e)$   
     $A \leftarrow getCluster(u)$   
     $B \leftarrow getCluster(v)$   
    if  $A \neq B$  then  
      Add edge  $e$  to  $T$   
      mergeClusters( $A, B$ )  
  return  $T$ 
```

Example



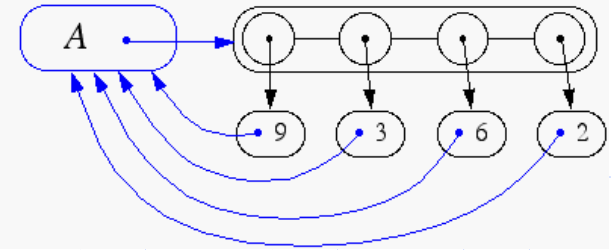
Example (contd.)



Data Structure for Kruskal's Algorithm

- ❑ The algorithm maintains a forest of trees
- ❑ A priority queue extracts the edges by increasing weight
- ❑ An edge is accepted if it connects distinct trees
- ❑ We need a data structure that maintains a **partition**, i.e., a collection of disjoint sets, with operations:
 - **makeSet**(u): create a set consisting of u
 - **find**(u): return the set storing u
 - **union**(A, B): replace sets A and B with their union

Recall of List-based Partition



- Each set is stored in a sequence
- Each element has a reference back to the set
 - operation **find**(u) takes $O(1)$ time, and returns the set of which u is a member.
 - in operation **union**(A,B), we move the elements of the smaller set to the sequence of the larger set and update their references
 - the time for operation **union**(A,B) is $\min(|A|, |B|)$
- Whenever an element is processed, it goes into a set of size at least double, hence each element is processed at most $\log n$ times

Partition-Based Implementation

- Partition-based version of Kruskal's Algorithm
 - Cluster merges as unions
 - Cluster locations as finds
- Running time
 - $O((n + m) \log n)$
 - PQ operations
 - $O(m \log n)$
 - UF operations
 - $O(n \log n)$

Algorithm *KruskalMST*(G)

Initialize a partition P

for each vertex v in G do

$P.makeSet(v)$

let Q be a priority queue.

Insert all edges into Q

$T \leftarrow \emptyset$

{ T is the union of the MSTs of the clusters}

while T has fewer than $n - 1$ edges do

$e \leftarrow Q.removeMin().getValue()$

$[u, v] \leftarrow G.endVertices(e)$

$A \leftarrow P.find(u)$

$B \leftarrow P.find(v)$

if $A \neq B$ then

Add edge e to T

$P.union(A, B)$

return T

Prim-Jarnik's Algorithm

- Similar to Dijkstra's algorithm
- We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s
- We store with each vertex v label $d(v)$ representing the smallest weight of an edge connecting v to a vertex in the cloud
- At each step:
 - We add to the cloud the vertex u outside the cloud with the smallest distance label
 - We update the labels of the vertices adjacent to u

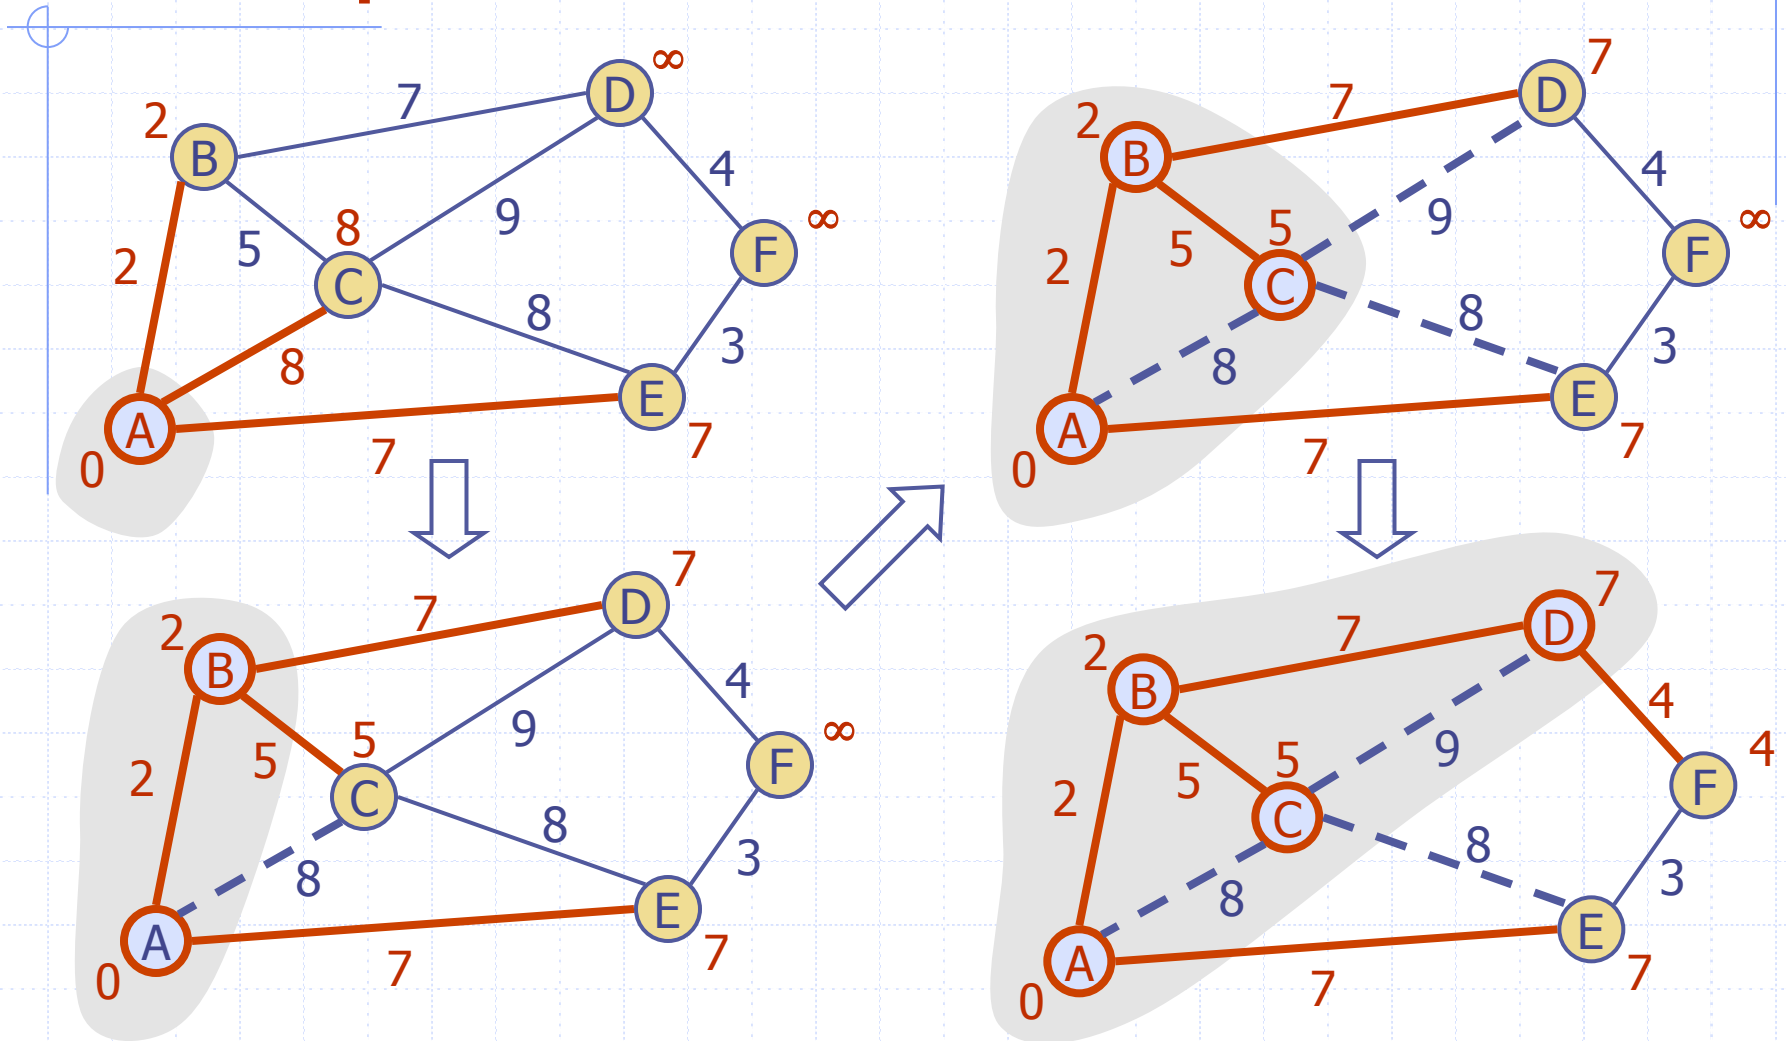
Prim-Jarnik's Algorithm (cont.)

- A heap-based adaptable priority queue with location-aware entries stores the vertices outside the cloud
 - Key: distance
 - Value: vertex
 - Recall that method *replaceKey(l,k)* changes the key of entry *l*
- We store three labels with each vertex:
 - Distance
 - Parent edge in MST
 - Entry in priority queue

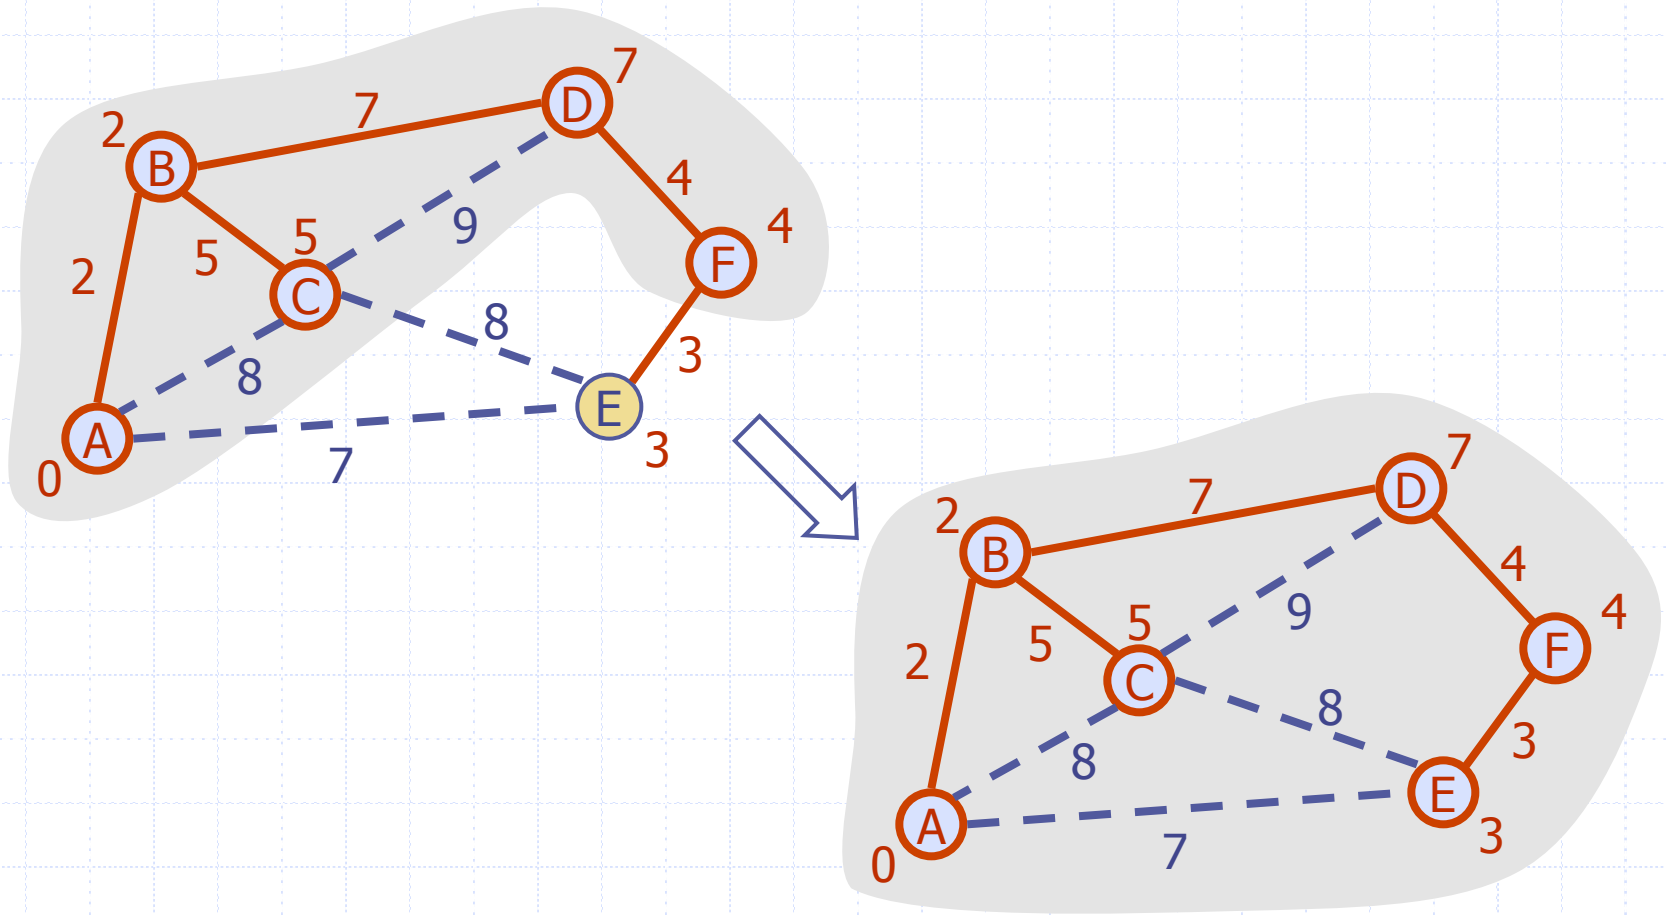
Algorithm *PrimJarnikMST(G)*

```
Q ← new heap-based priority queue
s ← a vertex of G
for all v ∈ G.vertices()
    if v = s
        v.setDistance(0)
    else
        v.setDistance(∞)
        v.setParent(∅)
        l ← Q.insert(v.getDistance(), v)
        v.setLocator(l)
while ¬Q.empty()
    l ← Q.removeMin()
    u ← l.getValue()
    for all e ∈ u.incidentEdges()
        z ← e.opposite(u)
        r ← e.weight()
        if r < z.getDistance()
            z.setDistance(r)
            z.setParent(e)
            Q.replaceKey(z.getEntry(), r)
```

Example



Example (contd.)



Analysis

- Graph operations
 - Method `incidentEdges` is called once for each vertex
- Label operations
 - We set/get the distance, parent and locator labels of vertex z $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex w in the priority queue is modified at most $\deg(w)$ times, where each key change takes $O(\log n)$ time
- Prim-Jarnik's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time is $O(m \log n)$ since the graph is connected

Baruvka's Algorithm (Exercise)

- Like Kruskal's Algorithm, Baruvka's algorithm grows many clusters at once and maintains a forest T
- Each iteration of the while loop halves the number of connected components in forest T
- The running time is $O(m \log n)$

Algorithm *BaruvkaMST*(G)

```
 $T \leftarrow V$  {just the vertices of  $G$ }  
while  $T$  has fewer than  $n - 1$  edges do  
  for each connected component  $C$  in  $T$  do  
    Let edge  $e$  be the smallest-weight edge from  $C$  to another component in  $T$   
    if  $e$  is not already in  $T$  then  
      Add edge  $e$  to  $T$   
return  $T$ 
```


Example of Baruvka's Algorithm (animated)

Slide by Matt Stallmann
included with permission.

